

Ben Bubnick

Senior Healthcare Data Consultant

Email: ben.bubnick@gmail.com | Phone: (216) 556-1612

LinkedIn: linkedin.com/in/benbubnick | GitHub: github.com/bubnickb

Professional Summary

Senior Healthcare Data Consultant with 10+ years designing ETL and ELT pipelines, cloud data platforms, analytics models, and data quality frameworks across clinical, payer, revenue cycle, and operational domains. Hands on experience with SQL, Python, Snowflake, dbt, SQLMesh, AWS, Azure, Spark, and healthcare interoperability standards including HL7 and FHIR. Leads requirements gathering, source to target mapping, data governance, reconciliation, stakeholder delivery, and technical documentation in HIPAA regulated environments.

Technical Skills

Data Engineering: SQL, Python, Snowflake, dbt, SQLMesh, Spark, Hadoop/HDFS, Hive, Impala, Kafka, ETL, ELT, data integration, data warehouses, data lakes

Cloud and DevOps: AWS (S3, Athena, Redshift, Glue), Azure (Data Factory, Blob Storage, Functions, SQL Database, Synapse Analytics), Docker, Git/GitHub, GitHub Actions, Jenkins, Azure DevOps

Healthcare Data: Revenue cycle, payer and eligibility data, clinical and financial data, HL7 v2/v3, FHIR, EDI/837, HEDIS, EHR/EMR, HIPAA, HITRUST

Analytics and BI: pandas, Jupyter, Power BI, Tableau, reconciliation, root cause analysis, statistical modeling, data visualization

Governance and Delivery: Data quality, source to target mapping, metadata and lineage, business requirements, acceptance criteria, QA, audit controls, Jira, Confluence, stakeholder engagement

Professional Experience

HHAEExchange

Oct. 2025 - Jun. 2026

Healthcare Data Integration Consultant - Contract via Veersa Technologies

Technologies: Snowflake, dbt, SQL/Jinja, Python, Snowpark, pandas, Liquibase, GitHub Actions, FHIR/Aidbox

Led healthcare source analysis, data quality investigation, and Snowflake modeling to produce trusted, standardized datasets for reporting and downstream use.

- Engineered unified and analytics Snowflake/dbt models spanning 55 tables across two organizations and two source system configurations, reducing client onboarding time by 67%.
- Implemented reusable SQL/Jinja transformation logic for lifecycle states, null handling, sentinel dates, and source system variations across organization, office, caregiver, and patient domains.
- Created or materially redesigned 30 automated SQL/dbt data quality tests covering nulls, duplicates, referential integrity, business rules, lifecycle logic, and source to target reconciliation.
- Translated ambiguous business requirements into documented mappings, acceptance criteria, onboarding standards, and 12 data quality investigation runbooks.

Adonis

Mar. 2025 - Jul. 2025

Healthcare Revenue Cycle Data Consultant - Contract via Health Data Movers

Technologies: AWS S3, Snowflake, SQLMesh, SQL, Python, Docker

Translated payer and revenue cycle source data into standardized Snowflake models for repeatable client onboarding, analytics, and reporting.

- Built, documented, validated, and maintained 54 SQLMesh models supporting revenue-cycle analytics for a provider serving 4,598 families across 113 locations and processing 86,388 annual authorizations.
- Implemented automated row count, uniqueness, and not null audits across 54 models, achieving a 100% pass rate across 117+ configured audit applications at final delivery.
- Diagnosed undocumented source key versioning and implemented composite business grains with SHA-256 surrogate keys across 10 charge and payment models, eliminating duplicate key failures.
- Authored implementation playbooks, technical handoff documentation, and a QA runbook for a repeatable healthcare data validation workflow.

GrowData Analytics

Sep. 2024 - Oct. 2025

Independent Data Automation and Analytics Consultant

Technologies: Ruby, RSpec, OCR, LLM APIs, audit logging, automated testing

- Built and validated a preproduction document to data MVP that converted semi structured invoices into auditable transaction records, achieving 95% field level accuracy across a stratified 500 invoice test suite.
- Instrumented the MVP with automated tests, audit logging, retry controls, and confidence routing, achieving 88% straight through processing and under 1% system failure.

Arkos Health

Jun. 2023 - Jan. 2025

Senior Data Analyst

Technologies: AWS S3, Athena, Redshift, Glue, dbt, PySpark, Docker, Jupyter, Power BI, Tableau

- Architected a greenfield AWS data lake integrating 18 source systems through approximately 200 scheduled data assets and 200 dbt models, supporting six analysts across six state markets.
- Integrated Jupyter based analytics and authored platform, architecture, and operating documentation, enabling six analysts to independently use the governed data lake for recurring, executive, and BI reporting.

Merative

Jul. 2022 - Mar. 2023

Senior Data Scientist

Technologies: Azure Data Factory, Blob Storage, Functions, SQL Database, Synapse Analytics, Python, SQL, Azure DevOps

- Engineered a seven layer Python and SQL migration acceptance suite covering partition reconciliation, schema compatibility, null handling, duplicates, field level comparisons, aggregates, and business rules, achieving exact HDFS to Azure fidelity.
- Led HIPAA compliant deidentification validation across demographic, encounter, and coded clinical data representing approximately 300 million patient lives, achieving a 100% pass rate before sign off.

IBM

Jun. 2015 - Jul. 2022

Advanced through five roles delivering healthcare analytics, data architecture, client integration, data quality, and large scale migration validation.

Technologies: Spark, Hadoop/HDFS, Hive, Impala, Kafka, Python, SQL, PostgreSQL, Docker, Jenkins, Tableau, HL7

Senior Data Scientist

Nov. 2021 - Jul. 2022

- Developed a surgical scheduling optimizer and CNN/RNN case card prediction proof of concept, automating case card generation and resource allocation across operating rooms, time slots, and personnel.
- Engineered a Spark, Hive, and Impala analytics framework powering 50+ prior authorization lag time and approval bottleneck reports with daily operational and periodic regulatory and financial snapshots.

Senior Data Architect

Dec. 2020 - Nov. 2021

- Architected Spark and Hadoop pipelines integrating more than 100 billion structured and semi structured clinical records, supporting four EPM applications, 30+ risk models, and hundreds of measures and registries.
- Replaced static mappings with metadata driven instruction sets across approximately 200 source integrations, standardizing EDI batch schemas while enabling Python/Jupyter lineage tracing and end to end validation gates.

Data Architect

Nov. 2017 - Dec. 2020

- Designed SQL validation frameworks and client review gates across approximately 100 integrations, reducing Jira tracked validation rework by 25% and mapping rework by 45%.
- Led a multi organization healthcare integration spanning 70+ EHR sources, implementing provider scoped patient filtering to enforce data sharing rules and coordinating delivery with engineering teams.

Software Developer, Data Platform

Sep. 2016 - Nov. 2017

- Engineered an HL7 v2/v3 parsing pipeline in Pig for approximately 30 internal users, replacing a Mirth based workflow and eliminating \$30,000 in annual vendor costs.
- Built a Python, Ruby, and SQL extraction tool that reduced total client integration effort by 28%, measured through eliminated story points and lower average engineering hours per implementation.

Data Scientist

Nov. 2015 - Sep. 2016

- Directed technical delivery across 36 client data integrations, coordinating 12 remote contractors and establishing code review, QA, and issue resolution standards.
- Authored 125+ technical knowledge base articles, approximately 30% of the team's published content, to standardize integration development and troubleshooting practices.

Explorys

Jun. 2015 - Nov. 2015

Data Science Consultant - Contract via Randstad USA

Technologies: Hadoop/HDFS, Ruby, Pig, React, Node.js, JavaScript, SQL, Oracle, Jenkins

Supported clinical data integration and population analytics across more than 10 client data warehouses.

- Diagnosed Hadoop and Oracle production ingestion failures and implemented Jenkins backed scheduling and recovery controls across 10+ client data warehouses.

HSEA - ETL Brigade Volunteer

Feb. 2025 - May 2025

Built a reusable PostgreSQL model for community health analytics covering claims, visits, and eligibility.

Education and Credentials

Master of Science, Astrophysics, University of Cincinnati, 2010

Certificate of Specialization, Data Science, Johns Hopkins University, 2015

IBM Data Science Professional Certificate, IBM, 2022

Bachelor of Fine Arts, Printmaking, Ohio University, 2004