

Ben Bubnick

HEALTHCARE DATA STRATEGY · DATA GOVERNANCE · ANALYTICS DELIVERY

✉ ben.bubnick@gmail.com | 📱 bubnicbf | 🌐 benbubnick

(216) 556-1612

Summary

Healthcare Data Consultant with 10+ years delivering governed analytics platforms and ML solutions across clinical, operational, and revenue-cycle domains. Combines hands-on Python, SQL, Snowflake, AWS, and Azure expertise with leadership in data quality, reconciliation, technical standards, and cross-functional delivery.

Experience

HHaEXchange

HEALTHCARE DATA INTEGRATION CONSULTANT — CONTRACT VIA VEERSA TECHNOLOGIES

Oct. 2025 - Jun. 2026

Led healthcare source analysis, data-quality investigation and Snowflake modeling to produce trusted, standardized datasets for reporting and downstream use.

- Engineered unified and analytics Snowflake/dbt models spanning 55 tables across two organizations and two source-system configurations, reducing client onboarding time by 67%.
- Implemented reusable SQL/Jinja logic for lifecycle states, null handling, sentinel dates, and source-system variations across organization, office, caregiver, and patient domains, centralizing 37% of post-source transformation logic in shared code.
- Created or materially redesigned 30 SQL/dbt data-quality tests and designed macros to automate null, duplicate, referential-integrity, business-rule, lifecycle, and source-to-target checks.
- Profiled source schemas and converted ambiguous business rules into documented mappings, acceptance criteria, and onboarding standards, plus 12 data-quality runbooks that standardized investigation and onboarding practices.

Adonis

HEALTHCARE REVENUE CYCLE DATA CONSULTANT — CONTRACT VIA HEALTH DATA MOVERS

Mar. 2025 - Jul. 2025

Translated payer and revenue-cycle source data into standardized Snowflake models for repeatable client onboarding and reporting.

- Built, documented, validated, and maintained 54 SQLMesh models, supporting revenue-cycle analytics for a provider serving 4,598 families across 113 locations and processing 86,388 annual authorizations.
- Implemented automated validation across 54 SQLMesh models, including exact row-count audits, uniqueness checks, and not-null checks; achieved a 100% pass rate across 117+ configured audit applications at final delivery, including exact row-count reconciliation for 40 models, with no unresolved discrepancies.
- Diagnosed undocumented source-key versioning and implemented composite business grains with SHA-256 surrogate keys across 10 charge and payment models in one day, eliminating duplicate-key failures and directly unblocking final delivery.
- Authored implementation playbooks, plus internal handoff documentation and a QA runbook that established a repeatable healthcare-data validation workflow for an engineering team new to the domain.

GrowData Analytics

INDEPENDENT DATA AUTOMATION & ANALYTICS CONSULTANT

Sep. 2024 - Oct. 2025

Built a document-to-data product that converted semi-structured invoices into validated, auditable transaction records.

- Built and validated a pre-production Ruby, OCR, and LLM invoice-automation MVP within 11 months, achieving 95% field-level accuracy across a stratified 500-invoice test suite (95% exact CI: 92.67%-96.79%).
- Instrumented the MVP with RSpec validation, audit logging, retry controls, error classification, and confidence-based routing, demonstrating 88% straight-through processing, 12% manual review, and under 1% complete system failure in testing.

Arkos Health

SENIOR DATA ANALYST

Jun. 2023 - Jan. 2025

Delivered healthcare analytics, recurring reporting and data-quality investigations while improving the governed data foundation used by the BI team.

- Architected a greenfield AWS data lake integrating 18 source systems through approximately 200 scheduled data assets and 200 dbt models, supporting six analysts across six state markets.
- Integrated Jupyter-based reporting and authored platform, architecture, and operating documentation, enabling six analysts to independently use the governed data lake for recurring and executive reporting.

Merative

SENIOR DATA SCIENTIST

Jul. 2022 - Mar. 2023

Led data validation, compliance and migration assurance for a regulated healthcare platform modernization.

- Engineered a seven-layer Python and SQL migration-acceptance suite covering partition-level reconciliation, schema compatibility, null handling, duplicates, field-level comparisons, aggregates, and business rules, achieving exact HDFS-to-Azure fidelity with zero unresolved discrepancies.
- Led HIPAA-compliant de-identification validation across demographic, encounter, and coded clinical data representing approximately 300 million patient lives, achieving a 100% pass rate before sign-off.

IBM

Advanced through five roles delivering healthcare analytics, client integration, data-quality improvement and large-scale migration validation.

SENIOR DATA SCIENTIST

Nov. 2021 - Jul. 2022

- Developed an MVP surgical-scheduling optimizer and a CNN/RNN case-card prediction proof of concept, automating case-card generation and resource allocation across operating rooms, time slots, and personnel.
- Engineered a reusable Spark, Hive, and Impala analytics framework powering 50+ downstream prior-authorization lag-time and approval-bottleneck reports, with daily operational execution and monthly or quarterly regulatory and financial snapshots.

SENIOR DATA ARCHITECT

Dec. 2020 - Nov. 2021

- Architected Spark and Hadoop pipelines that integrated and normalized more than 100 billion structured and semi-structured clinical records, supporting four EPM applications, 30+ risk models, and hundreds of standardized measures and registries.
- Replaced static mappings with executable, metadata-driven instruction sets adopted across approximately 200 source integrations, standardizing EDI batch schemas while enabling modular integration templates, automated Python/Jupyter lineage tracing, and end-to-end validation gates.

DATA ARCHITECT

Nov. 2017 - Dec. 2020

- Designed standardized SQL validation frameworks and client-review gates across approximately 100 client integrations, reducing Jira-tracked validation rework hours by 25% and mapping rework hours by 45% across initial delivery and subsequent corrections.
- Led architecture and delivery of a multi-organization healthcare integration spanning 70+ heterogeneous EHR sources and isolated organizational schemas; implemented provider-scoped patient filtering to enforce data-sharing rules, and coordinated with engineering teams.

EARLIER ROLES

Jun. 2015 - Nov. 2017

- Directed technical delivery across 36 client data integrations, coordinating 12 remote contractors and establishing code-review, QA, and issue-resolution standards.

Education & Credentials

2010 **Master of Science** Astrophysics
2015 **Certificate of Specialization** Data Science
2022 **IBM Data Science Professional Certificate**

University of Cincinnati
Johns Hopkins University
IBM